

Rohan Patel

Senior Backend Engineer (Distributed Systems, AI/LLM Systems)
LinkedIn: rohan-patel-5254aa19b

Email: rohanpatel98.rp@gmail.com
Mobile: +91 7700006792
GitHub: rohanpatel981

SUMMARY

Senior Backend Engineer with 5+ years of experience designing and scaling distributed systems across Amazon, Embibe / Reliance Jio (10M+ users), and Uniqus. Strong experience in high-throughput, low-latency services using Java, Spring Boot, TypeScript, and AWS, with deep work in concurrency, distributed locking, event-driven architecture, and horizontal scaling. Delivered a **60× API latency cut** and a **6–9× end-to-end workflow speedup** on a revenue-critical platform. Actively building production AI systems — multi-agent Claude pipelines, MCP-based tool integrations, RAG, and vector search. Daily Claude Code + MCP user. Twice awarded Best Performance Engineer.

EXPERIENCE

- **Uniqus Consultech** Mumbai, India
Apr 2024 - Present
Software Engineer II (Backend Tech Lead)
 - **System Ownership:** Architected and shipped a revenue-critical distributed platform spanning 15+ modules and multiple microservices; owned HLD/LLD, scaling, and production stability end-to-end.
 - **60× Latency Win:** Cut critical API latency from ~60s to under 1s via query optimization, caching, and bulk read/write on large-scale datasets.
 - **Workflow Throughput:** Drove end-to-end workflow runtime from ~4.5 minutes to 30–45 seconds (6–9×) by tuning pod count, infra sizing, and DB instance topology.
 - **Horizontal Scaling:** Migrated services from single-pod to multi-pod (3–4 pods in parallel) using Redis-based distributed locking and dirty-key marking for concurrency safety on shared state.
 - **Concurrency & Throughput:** Multi-pod rollout unlocked parallel request handling under high-concurrency workloads, lifting throughput across all flows and improving production reliability.
 - **Event-Driven Architecture:** Led RabbitMQ integration across services and AWS Lambdas to enable async processing, decoupling, and fault tolerance; integrated AWS SES for high-volume notifications.
 - **Reliability & Automation:** Designed an idempotent auto-update workflow that propagates changes across downstream modules with zero manual intervention, reducing operational overhead.
 - **Technical Leadership:** Lead a team of 5 engineers — code reviews, release approvals, on-call ownership; conducted 30+ SDE-I/II interviews and mentored juniors.
 - **Product Impact:** Improved system responsiveness and automation for a revenue-critical platform, directly enhancing end-user experience and reducing manual intervention.
- **Embibe** Bengaluru, India
Jan 2023 - May 2024
Software Engineer II
 - **High-Scale Backend:** Owned backend services for Reliance Jio's B2B education platform serving ~10M users; delivered state management, profile, and multi-profile switching APIs.
 - **Performance:** Optimized hot code paths and DB access patterns, improving throughput and latency by 30–40% on high-traffic workflows.
- **Amazon** Bangalore, India
Oct 2021 - Dec 2022
Software Development Engineer I
 - **Tier-1 Service Development:** Delivered 6 critical features in a Tier-1 checkout service powering 10M+ users with zero customer-facing downtime.
 - **Reliability & Observability:** Implemented Service Level Indicators and led a CASE campaign that reduced unavailability and throttling exceptions on a high-throughput service.
 - **Privacy & Compliance:** Implemented DSAR and on-demand data deletion across services in partnership with cross-functional teams.
 - **Service Ownership:** Acted as Tier-1 Service Owner with SDMs/TPMs; resolved recurring operational issues and significantly raised unit-test coverage on a Tier-2 service.

AI / LLM SYSTEMS EXPERIENCE

- **Multi-Agent Systems:** Designing and building production multi-agent pipelines using LLMs (Claude) to automate complex finance workflows, with orchestration, tool use, and structured handoffs between agents.
- **RAG Pipelines:** Built retrieval-augmented generation pipelines over large internal document corpora, integrating chunking, embeddings, and re-ranking for grounded, low-hallucination answers.
- **Vector Search:** Used vector databases for semantic search and context retrieval, indexing embeddings of domain documents to power agent reasoning at low query latency.
- **Document Ingestion:** Built ingestion pipelines for parsing, normalizing, chunking, and embedding heterogeneous documents into the retrieval store with idempotent, async processing.
- **Prompt & Context Engineering:** Designed prompt templates, system instructions, and context-window strategies to keep agent behavior reliable, deterministic, and cost-efficient under production load.
- **MCP & Agentic Tooling:** Daily user of MCP (Model Context Protocol) and Claude Code; build and consume MCP servers exposing internal tools to agents, with attention to OAuth-style auth, idempotent tool calls, and per-tool observability.

EDUCATION

- **University of Mumbai** Maharashtra, India
Bachelor of Engineering in Computer Engineering; CGPA: 7.65 *Aug 2016 - Oct 2020*
 - **Relevant Coursework:** Operating Systems, Algorithms, Data Structures, Computer Networks, Distributed Systems.

AWARDS & RECOGNITION

- **Best Performance Engineer — Uniqus Consultech:** Awarded twice in 2 years (2024, 2025) for sustained impact on system scalability, delivery, and technical leadership.

TECHNICAL SKILLS

- **Languages:** Java, TypeScript, Python, C++, SQL, C#
- **Backend & Frameworks:** Spring Boot, Spring Framework, Spring Cloud Data Flow, REST APIs, Microservices
- **Distributed Systems:** HLD/LLD, Horizontal Scaling, Distributed Locking, Concurrency, Idempotency, Fault Tolerance, Caching, High-Throughput & Low-Latency Systems, Parallel Processing
- **Messaging & Event-Driven:** Kafka, RabbitMQ, AWS SQS, Async Processing, Event-Driven Architecture
- **Databases & Caching:** MongoDB, DynamoDB, MySQL, Redis, Vector Databases
- **Cloud & DevOps:** AWS (Lambda, SQS, SES), Azure, Jenkins, Git, GitHub, GitLab, Bitbucket
- **AI Systems / LLM:** LLM Integration (Claude), Multi-Agent Systems, RAG, Vector Search, Embeddings, Prompt Engineering, Context Engineering, MCP (Model Context Protocol), Agentic Workflows, Claude Code
- **CS Fundamentals:** Data Structures, Algorithms, OOP, Operating Systems, Computer Networks
- **Leadership:** Technical leadership, mentorship, code reviews, hiring (30+ interviews), cross-functional collaboration